

coursework_01

January 27, 2026

1 Coursework 1: Image filtering

The coursework includes coding questions and/or written questions. Please read both the text and the code in this notebook to get an idea what you are expected to implement.

1.1 What is expected?

- Complete and run the code using `jupyter-lab`.
- Export (File | Save and Export Notebook As...) the notebook as a PDF file, which contains your code, results and answers, and upload the PDF file onto [Scientia](#).
- Instead of clicking the Export button, you can also run the following command instead:
`jupyter nbconvert coursework_01.ipynb --to pdf`
- If Jupyter complains issues during exporting, it is likely that `pandoc` or `latex` is not installed, or their paths have not been included. You can install the relevant libraries and retry. Alternatively, use the Print function of your browser to export the PDF file.
- If Jupyter-lab does not work for you at the end, alternatively, you can use Google Colab to write the code and export the PDF file.

1.2 Dependencies:

You may need to install [Jupyter-Lab](#) and other libraries used in this coursework, such as by running the command: `pip3 install [package_name]`

```
[1]: # Import libraries (provided)
import imageio.v3 as imageio
import numpy as np
import matplotlib.pyplot as plt
import scipy
import scipy.signal
import math
import time
```

1.3 Q1. Moving average filtering (20 points).

Read the provided clean image, add noise to the image and design a moving average filter for denoising.


You are expected to design the kernel of the filter and then perform 2D image filtering using the function `scipy.signal.convolve2d()`.

```
[2]: # Read the image (provided)
image = imageio.imread('campus_snow.jpg')

# Corrupt the image with Gaussian noise (provided)
noise_mu = 0
noise_sigma = 50
noise = np.random.normal(noise_mu, noise_sigma, image.shape)
image_noisy = image + noise

# Visualise the images (provided)
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Clean image')

plt.subplot(1, 2, 2)
plt.imshow(image_noisy, cmap='gray')
plt.title('Noisy image')
plt.gcf().set_size_inches(8, 4)
```



1.3.1 Q1.1 Filter the noisy image using a 5x5 moving average filter. Display the filtered image.

```
[ ]: # Design the filter h
h = np.full((5, 5), 0.04)
```

```


# Convolve the noisy image with h using scipy.signal.convolve2d function
image_filtered = scipy.signal.convolve2d(image_noisy, h, mode='same')

# Print the filter (provided)
print('Filter h = {0}'.format(h))

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('After moving average filtering')
plt.gcf().set_size_inches(4, 4)

```

Filter h = [[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]
[0.04 0.04 0.04 0.04 0.04]]



1.3.2 Q1.2 Assess the quality of the filtered image using a quantitative metric, the peak signal-to-noise ratio (PSNR).

For this case, the pixel intensity of the image is represented using the uint8 format, with the peak value to be 255. The PSNR is defined as,

$$\text{PSNR} = 10 \cdot \log_{10} \frac{255^2}{\frac{1}{N} \sum_x [J(x) - I(x)]^2}$$

where x denotes the pixel index, N denotes the total number of pixels in the image, J denotes the

filtered i.e. denoised image and I denotes the ground truth clean image. The denominator of the term within the logarithm operator is the mean squared error between I and J .

You can find more detail about PSNR [here](#).

```
[ ]: # Implement the PSNR function
def eval_psnr(I, J):
    # I: the ground truth clean image (peak value: 255 for uint8 data format)
    # J: the denoised image
    #
    # return: the PSNR metric (unit: dB)

    N = I.shape[0] * I.shape[1]
    mse = np.sum((J - I) ** 2) / N
    psnr = 10 * math.log10(255 ** 2 / mse)
    return psnr

# Evaluate the PSNR for the filtered image (provided)
psnr = eval_psnr(image, image_filtered)

# Print the PSNR (provided)
print('PSNR = {:.2f} dB'.format(psnr))
```

PSNR = 18.39 dB

1.4 Q2. Gaussian filtering (70 points).

1.4.1 Q2.1 Implement a function that constructs a 2D Gaussian filter given the parameter σ .


```
[ ]: # Implement the Gaussian filter
def gaussian_filter_2d(sigma):
    # sigma: the parameter sigma for the Gaussian kernel (unit: pixel)
    #
    # return: a 2D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))

    # Calculate the filter weights
    x, y = np.meshgrid(np.arange(-rad, rad+1), np.arange(-rad, rad+1))
    h = np.exp(-(x**2 + y**2) / (2 * sigma**2)) / (2 * math.pi * sigma**2)
    h = h / np.sum(h)
    return h

# Visualise the Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)
plt.imshow(h)
```

```
plt.title('Gaussian kernel')
plt.gcf().set_size_inches(4, 4)
```




1.4.2 Q2.2 Perform Gaussian filtering ($\sigma = 5$ pixels) for the noisy image, evaluate the computational time for Gaussian filtering and display the filtered image.

```
[ ]: # Construct the Gaussian filter (provided)
sigma = 5
h = gaussian_filter_2d(sigma)

# Perform Gaussian filtering and count time
start_time = time.time()
image_smoothed_2d = scipy.signal.convolve2d(image_noisy, h, mode='same')
time_elapsed = time.time() - start_time
print("Time elapsed in seconds: {:.3f}".format(time_elapsed))

# Visualise the filtered image (provided)
plt.imshow(image_smoothed_2d, cmap='gray')
plt.title('Gaussian smoothed image')
plt.gcf().set_size_inches(4, 4)
```

Time elapsed in seconds: 4.238



1.4.3 Q2.3 Implement a function that generates a 1D Gaussian filter given the parameter σ . Construct 1D Gaussian filters along x-axis and y-axis respectively.

```
[ ]: # Implement the 1D Gaussian filter
def gaussian_filter_1d(sigma):
    # sigma: the parameter sigma in the Gaussian kernel (unit: pixel)
    #
    # return: a 1D array for the Gaussian kernel

    # The filter radius is 4 times sigma (provided)
    rad = int(math.ceil(4 * sigma))

    # Calculate the filter weights
    h = np.arange(-rad, rad + 1)
    h = np.exp(-(h**2) / (2 * sigma**2)) / (2 * math.pi * sigma**2)
    h = h / np.sum(h)
    return h

# sigma = 5 pixel (provided)
sigma = 5

# Construct the Gaussian filter along x-axis. Its shape is (1, sz).
h_x = gaussian_filter_1d(sigma).reshape(1, -1)
```


```

# Construct the Gaussian filter along y-axis. Its shape is (sz, 1).
h_y = gaussian_filter_1d(sigma).reshape(-1, 1)

# Visualise the filters (provided)
plt.subplot(1, 2, 1)
plt.imshow(h_x)

plt.subplot(1, 2, 2)
plt.imshow(h_y)
plt.gcf().set_size_inches(4, 4)

```



1.4.4 Q2.4 Perform Gaussian filtering ($\sigma = 5$ pixels) using two separable filters and evaluate the computational time for separable Gaussian filtering. Compare the smoothed image using separable filtering to the smoothed image using a single 2D Gaussian filter.

```
[ ]: # Perform separable Gaussian smoothing and count time
start_time = time.time()
horizontal_filtered_image = scipy.signal.convolve2d(image_noisy, h_x, mode='same')
image_smoothed = scipy.signal.convolve2d(horizontal_filtered_image, h_y, mode='same')
time_elapsed = time.time() - start_time
print("Time elapsed in seconds: {0:.3f}".format(time_elapsed))
```

```
# Report the difference between the separably filtered image and the image
# filtered by a single 2D Gaussian filter (provided)
diff = image_smoothed - image_smoothed_2d
print('Mean absolute difference = {:.6f}'.format(np.mean(np.abs(diff))))
```

Time elapsed in seconds: 0.531
 Mean absolute difference = 0.000000


1.4.5 Q2.5 Perform Gaussian smoothing for the same noisy image, assess the quality of the Gaussian smoothed image using PSNR, when different sigma values are used.

```
[ ]: # A list of sigma values (provided)
list_sigma = np.arange(0.5, 5, 0.5)

# Perform Gaussian smoothing with different sigma values and record the PSNR
# values
list_psnr = []
for sigma in list_sigma:
    h_x = gaussian_filter_1d(sigma).reshape(1, -1)
    h_y = gaussian_filter_1d(sigma).reshape(-1, 1)
    horizontal_filtered_image = scipy.signal.convolve2d(image_noisy, h_x,
    mode='same')
    image_smoothed = scipy.signal.convolve2d(horizontal_filtered_image, h_y,
    mode='same')
    psnr = eval_psnr(image, image_smoothed)
    list_psnr.append(psnr)

# Plot the PSNR metric against sigma (provided)
plt.plot(list_sigma, list_psnr)
plt.xlim([0, 5])
plt.xlabel('$\sigma$ (unit: pixel)')
plt.ylabel('PSNR (unit: dB)')
plt.grid()
plt.gcf().set_size_inches(4, 4)
```

<>:19: SyntaxWarning: invalid escape sequence '\s'
 <>:19: SyntaxWarning: invalid escape sequence '\s'
 /tmp/ipykernel_3424195/2877778068.py:19: SyntaxWarning: invalid escape sequence '\s'
 plt.xlabel('\$\sigma\$ (unit: pixel)')



1.4.6 Q2.6 Implement 3x3 Sobel filters, perform Sobel filtering for the noisy image, and display the gradient magnitude map.

```
[17]: # Construct the Sobel filters
sobel_x = np.array([[1, 0, -1], [2, 0, -2], [1, 0, -1]])
sobel_y = np.array([[1, 2, 1], [0, 0, 0], [-1, -2, -1]])

# Print the filters (provided)
print('Sobel_x = {}'.format(sobel_x))
print('Sobel_y = {}'.format(sobel_y))

# Sobel filtering for the noisy image
grad_x = scipy.signal.convolve2d(image_noisy, sobel_x, mode='same')
grad_y = scipy.signal.convolve2d(image_noisy, sobel_y, mode='same')


# Calculate the gradient magnitude
### Insert your code ####
grad_mag_noisy = np.sqrt(grad_x**2 + grad_y**2)

# Display the magnitude map (provided)
plt.imshow(grad_mag_noisy, cmap='gray', vmin=0, vmax=500)
plt.title('Gradient magnitude map')
plt.gcf().set_size_inches(4, 4)
```

```

Sobel_x = [[ 1  0 -1]
           [ 2  0 -2]
           [ 1  0 -1]]
Sobel_y = [[ 1  2  1]
           [ 0  0  0]
           [-1 -2 -1]]

```



1.4.7 Q2.7 Perform Gaussian smoothing for the noisy image, followed by Sobel filtering and display the gradient magnitude map.

```
[18]: # Parameter for the Gaussian filter (provided)
sigma = 5

# Gaussian smoothing
gaussian_x = gaussian_filter_1d(sigma).reshape(1, -1)
gaussian_y = gaussian_filter_1d(sigma).reshape(-1, 1)
horizontal_filtered_image = scipy.signal.convolve2d(image_noisy, gaussian_x, mode='same')
image_smoothed = scipy.signal.convolve2d(horizontal_filtered_image, gaussian_y, mode='same')

# Sobel filtering
grad_x = scipy.signal.convolve2d(image_smoothed, sobel_x, mode='same')
grad_y = scipy.signal.convolve2d(image_smoothed, sobel_y, mode='same')
```


```

# Calculate the gradient magnitude
### Insert your code ####
grad_mag = np.sqrt(grad_x**2 + grad_y**2)

# Display the magnitude map (provided)
plt.imshow(grad_mag, cmap='gray', vmin=0, vmax=100)
plt.title('Gradient magnitude map after Gaussian smoothing')
plt.gcf().set_size_inches(4, 4)

```

Gradient magnitude map after Gaussian smoothing



1.5 Q3. Implement image filters using Pytorch (10 points).

Pytorch is a machine learning framework that supports filtering and convolution.

The `Conv2D` operator takes an input array of dimension $N \times C_1 \times X \times Y$, applies the filter and outputs an array of dimension $N \times C_2 \times X \times Y$. Here, since we only have one image with one colour channel, we will set $N=1$, $C_1=1$ and $C_2=1$. You can read the documentation of `Conv2D` for more detail.

```
[20]: # Import libraries (provided)
import torch
```

1.5.1 Q3.1 Expand the dimension of the noisy image into 1x1xXxY and convert it to a Pytorch tensor.

```
[23]: # Expand the dimension of the numpy array
image_noisy_expanded = image_noisy[np.newaxis, np.newaxis, :, :]

# Convert to a Pytorch tensor using torch.from_numpy
### Insert your code ####
tensor_noisy = torch.from_numpy(image_noisy_expanded).float()
```

1.5.2 Q3.2 Create a Pytorch Conv2D filter, set its kernel to be a 2D Gaussian filter, perform filtering, report computational time and display the result.

```
[26]: # A 2D Gaussian filter when sigma = 5 pixel (provided)
sigma = 5
h = gaussian_filter_2d(sigma)


# Construct the Conv2D filter
h_expanded = h[np.newaxis, np.newaxis, :, :]
h_tensor = torch.from_numpy(h_expanded).float()

# Filtering and assess computational time
start_time = time.time()
image_filtered_tensor = torch.nn.functional.conv2d(tensor_noisy, h_tensor, □
    ↪padding='same')
time_elapsed = time.time() - start_time
print("Time elapsed in seconds: {:.3f}".format(time_elapsed))
image_filtered = image_filtered_tensor.squeeze().detach().numpy()

# Display the filtering result (provided)
plt.imshow(image_filtered, cmap='gray')
plt.title('Gaussian smoothed image using PyTorch')
plt.gcf().set_size_inches(4, 4)
```

Time elapsed in seconds: 6.554

Gaussian smoothed image using PyTorch



[]: